



The Gaussian Classifier

Machine Learning – Laboratory


Battista Biggio

Department of Electrical and Electronic Engineering
University of Cagliari, Italy

Gaussian Classifier - Prediction

- Bayesian classifier that uses MAP for predictions: $y_k^* = \arg \max_{y_k} p(y_k|x) = \frac{p(x|y_k)p_k}{p(x)}$
- The likelihood is computed using a multivariate Gaussian distribution
 - This means that each class is modeled as Gaussian, with parameters μ_k, Σ_k


$$p(x|y_k) = g(x; \mu_k, \Sigma_k) = \frac{1}{\sqrt{(2\pi)^d \det \Sigma_k}} \exp\left(-\frac{1}{2}(x - \mu_k)^T \Sigma_k^{-1}(x - \mu_k)\right)$$



- Recall also that:
 - p_k is the prior probability of class y_k ;
 - the evidence is obtained, as usual, by marginalization over the classes, i.e., $p(x) = \sum_k p(x|y_k)p_k$

Gaussian Classifier - Training

- During training, we aim to fit one Gaussian distribution per class.
 - But how do we find the best parameters μ_k, Σ_k for each Gaussian?
 - What is the meaning of *best parameters* here?



Which of these Gaussian distributions is a better fit to the 'x' data points? Why?

Gaussian Classifier - Training

- **Maximum Likelihood Estimation (MLE).** MLE defines the best parameters of our model θ^* as those maximizing the likelihood L that the data x_1, \dots, x_n is generated by the model:

$$\theta^* = \underset{\theta}{\operatorname{argmax}} L(x_1, \dots, x_n | \theta)$$

- In our case, we aim to fit one Gaussian distribution per class. Thus, for each class,
 - we first extract the samples x_1, \dots, x_n for that class, and then
 - estimate the parameters $\theta = (\mu_k, \Sigma_k)$ for that class via MLE.
- **Good news:** MLE for Gaussian fitting can be solved in closed form!
 - The optimal μ_k^*, Σ_k^* values are obtained as the **sample mean** and the **sample covariance**
$$\mu_k^* = \frac{1}{n} \sum_{i=1}^n x_i$$
$$\Sigma_k^* = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu_k^*)(x_i - \mu_k^*)^T$$
 - ... and these equations are already implemented in many statistical tools and libraries!

Exercise 1

- Implement a Gaussian classifier that
 1. fits a Gaussian distribution to each training class
 2. predicts test samples by assigning them to the class with the max. a-posteriori probability (MAP)
- Test it on some Gaussian dataset
 - Measure test error
 - Visualize the decision regions

Solution (small excerpt from the notebook code)

```
class CClassifierGaussian:
    """
    Class implementing a Gaussian classifier
    """

    def fit(self, x, y):
        """Estimate priors, centroids and covariances with
        maximum likelihood estimates from the training data x,y.
        """
        n_classes = np.unique(y).size
        n_features = x.shape[1]

        self._priors = np.zeros(shape=(n_classes,))
        self._centroids = np.zeros(shape=(n_classes, n_features))
        self._covariances = np.zeros(shape=(n_classes, n_features, n_features))

        for k in range(n_classes):
            self._centroids[k, :] = x[y == k, :].mean(axis=0)
            self._priors[k] = (y == k).mean()
            self._covariances[k, :, :] = np.cov(x[y == k, :].T)

        self._priors /= self._priors.sum() # ensure priors sum up to 1
        return self
```

These are the maximum-likelihood estimates for the parameters of each class

Solution (small excerpt from the notebook code)

```
def decision_function(self, x):  
    """Return posterior estimates for each class"""  
  
    n_samples = x.shape[0]  
    n_classes = self._centroids.shape[0]  
    scores = np.zeros(shape=(n_samples, n_classes))
```

The matrix scores will contain, for each row (sample), the posterior probability predicted for each class.

Each sample will then be assigned to the class exhibiting the maximum a-posteriori probability estimate

class labels

0	1	2
0.1	0.2	0.7
0.3	0.1	0.6

Solution (small excerpt from the notebook code)

```
def decision_function(self, x):  
    """Return posterior or joint probability estimates for each class,  
    depending on whether posterior=True or False."""  
    n_samples = x.shape[0]  
    n_classes = self._centroids.shape[0]  
    scores = np.zeros(shape=(n_samples, n_classes))  
    for k in range(n_classes):  
        likelihood_k = mvn.pdf(  
            x, mean=self._centroids[k, :], cov=self._covariances[k, :, :])  
        scores[:, k] = self._priors[k] * likelihood_k # joint probability  
  
    if self.posterior:  
        # if posterior probs are required, divide joint probs by evidence  
        evidence = scores.sum(axis=1)  
        for k in range(n_classes):  
            # normalize per row to estimate posterior  
            scores[:, k] /= evidence  
    return scores  
  
def predict(self, x):  
    """Return predicted labels."""  
    scores = self.decision_function(x)  
    y_pred = np.argmax(scores, axis=1)  
    return y_pred
```

Compute Gaussian PDF at x.
This is our likelihood:

$$\frac{1}{\sqrt{(2\pi)^d \det \Sigma_k}} \exp\left(-\frac{1}{2}(x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k)\right)$$

Then multiply it by the prior to
obtain the joint probability, and
divide by the evidence $p(x)$ to
obtain the posterior

Take the index of the maximum
posterior to predict the class label

Solution

```
n_samples = [500, 500, 500]
centroids = [[-5, -5],
              [+5, -5],
              [0, +5]]
cov=[[ [3, -1],
        [-1, 3]],
      [ [3, -0.5],
        [-0.5, 3]],
      [ [1, -1],
        [-1, 3]]]


# generate data
x_tr, y_tr = make_gaussian_dataset(n_samples, centroids, cov=cov)
x_ts, y_ts = make_gaussian_dataset(n_samples, centroids, cov=cov)

clf = CClassifierGaussian()
clf.fit(x_tr, y_tr)
plot_decision_regions(x_tr, y_tr, classifier=clf)
plot_dataset(x_tr, y_tr)
plt.show()

scores = clf.decision_function(x_ts)
y_pred = clf.predict(x_ts)

print('Estimated priors: ', clf.priors)
print('Estimated centroids (with MLE): ', clf.centroids)
print('Estimated covariances (with MLE): ', clf.covariances)

print('Test error (%): ', (y_pred != y_ts).mean()*100)
```



Estimated priors:
[0.33333333 0.33333333 0.33333333]

Estimated centroids (with MLE):
[[-5.02818576 -5.07365035]
[4.95564731 -5.04938789]
[-0.03591899 4.87763586]]

Estimated covariances (with MLE):
[[[2.89255506 -1.05314784]
[-1.05314784 2.77306075]]

[[2.88868113 -0.46682182]
[-0.46682182 3.04256309]]

[[1.04041627 -1.09369177]
[-1.09369177 3.26337424]]]

Test error (%): 0.26666666666666666

Exercise 2


- Visualize the optimal decision regions for a 3-class Gaussian Classifier with parameters:
 - $p_0 = p_1 = \frac{1}{4}, p_2 = \frac{1}{2}$
 - $\mu_0 = [0, 0]^T, \mu_1 = [2, 2]^T, \mu_2 = [-2, -2]^T$
 - $\Sigma_0 = \Sigma_1 = \Sigma_2 = \begin{bmatrix} 4 & -3 \\ -3 & 4 \end{bmatrix}$

Solution

- In this case, we set the parameters directly to the Gaussian Classifier (instead of estimating them via *fit*), and then visualize the decision regions and boundaries

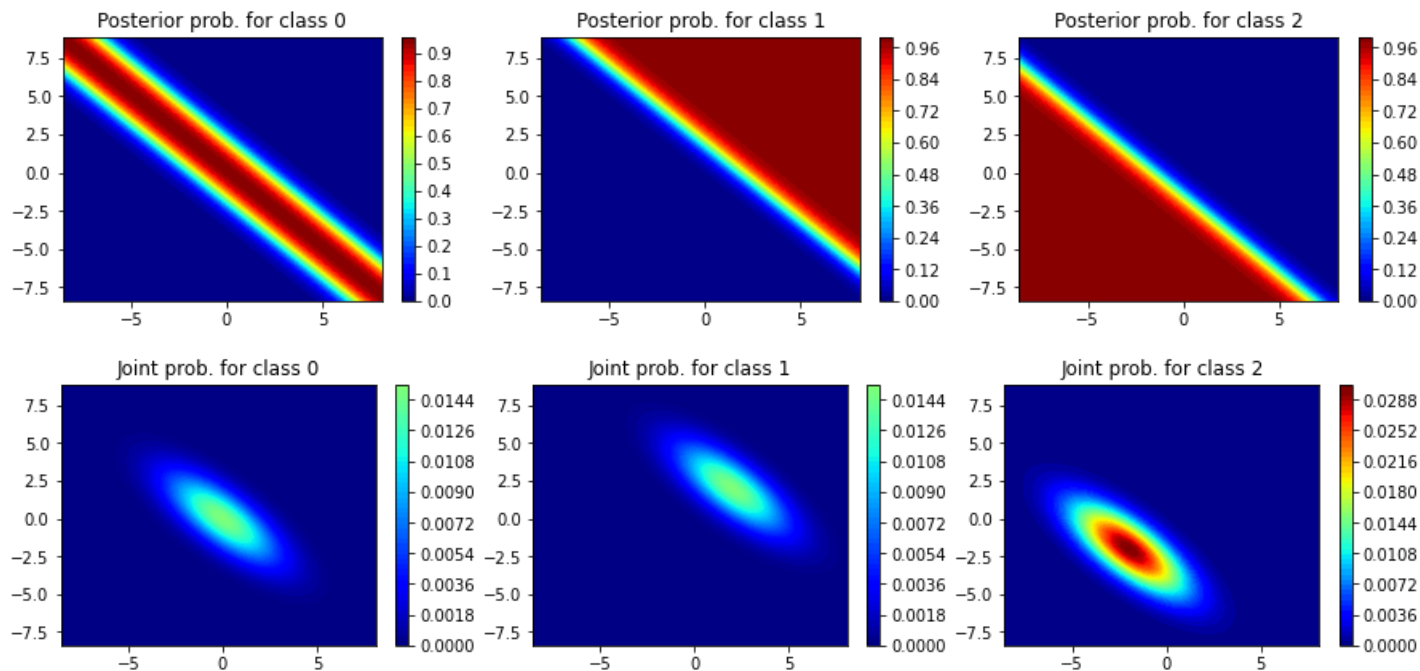
```
priors = np.array([1/4, 1/4, 1/2])
centroids = [[0, 0],
              [2, 2],
              [-2, -2]]
cov=[[4, -3], [-3, 4]],
     [[4, -3], [-3, 4]],
     [[4, -3], [-3, 4]]

clf = CClassifierGaussian()
# we do not estimate the parameters here,
# but use the true ones
# clf.fit(x_tr, y_tr)
clf._priors = np.array(priors)
clf._centroids = np.array(centroids)
clf._covariances = np.array(cov)
plot_decision_regions(x_tr, y_tr, classifier=clf)
plot_dataset(x_tr, y_tr)
plt.show()
```




Solution

- We can also visualize the posterior and joint probabilities for each class



Solution

- Finally, we also plot the evidence $p(x)$ along the z-axis in a 3D plot



- Note here the overlap among the three Gaussian distributions, one per class, and that the Gaussian associated to the highest prior (1/2) has a higher peak