

Part 3:

Elements of Non-parametric Techniques: the k-Nearest Neighbor (kNN) Classifier


Battista Biggio
battista.biggio@unica.it

Introduction

- In Part 2, we assumed that the forms of the probability density functions were known
 - However, this assumption cannot be made in some pattern recognition applications
- In this chapter, we examine a nonparametric method that can be used with arbitrary distributions, without assuming knowledge of the underlying probability densities
- We discuss the k-Nearest Neighbor (kNN) pattern classifier, which allows:
 - estimating the density function $p(\mathbf{x} | \omega_j)$
 - estimating the posterior probability $P(\omega_j | \mathbf{x})$

The k-Nearest Neighbor (kNN) Method

- Nonparametric classification is often associated with the notion of **prototype**
- We can think of a prototype as a representative element from a class
- The class label assigned to an example is based on the similarity of this example to one or more prototypes
- Similarity is defined in a geometrical sense, i.e., based on a certain distance
 - The smaller the distance, the higher the similarity between the input \mathbf{x} and the prototype(s)




The k-Nearest Neighbor (kNN) Method

- *kNN is one of the most theoretically elegant and simple classification techniques (L. Kuncheva, 2004)*
- Let D be a labeled training set containing n points, called prototypes
- Each prototype belongs to one of the “ c ” classes, i.e., $\mathbf{x}_i \in \omega_j \quad j=1, \dots, c$


$$\mathbf{D} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n]$$

$$\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{id}) \quad i=1, \dots, n$$



The k-Nearest Neighbor (kNN) Method

- To classify an input \mathbf{x} , the k nearest prototypes are retrieved from \mathbf{D} together with their class labels
- The input \mathbf{x} is labeled to the most represented class label amongst the k nearest neighbors
- In the figure:
 - \mathbf{x} is the pattern to be classified
 - we consider a “region” R of the feature space containing the k nearest prototypes to \mathbf{x}
 - we classify \mathbf{x} as belonging to the most represented class label amongst the k nearest neighbors within the region R .



The k-Nearest Neighbor (kNN) Method

- It can be shown that the kNN method estimates the posterior probabilities as:

$$\hat{P}(\omega_i|\mathbf{x}) = \frac{k_i}{k}$$

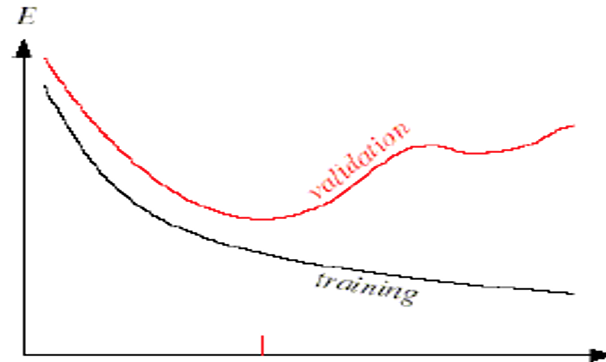
- k_i is the number of nearest neighbors belonging to the class ω_i within region R
- k is the number of the **k nearest prototypes** of \mathbf{x} within region R
- The minimum (Bayes) error classifier using the approximations above will assign \mathbf{x} to the class with the highest posterior probability
 - i.e., the class most represented among the k nearest neighbors of \mathbf{x}

How to Select the Right Value of the Hyperparameter k

- One simple rule of thumb (heuristic rule) is to select k as $k = \sqrt{n}$
 - n is the number of examples in our training set D
- Note that this rule imposes that the number k of nearest neighbors of the pattern \mathbf{x} that fall within the region R is smaller than the total number n of training set examples
- In binary (two-class) classification problems or problems with an **even** number of classes, it is helpful to choose k to be an odd number as this avoids **tie-breaks**

How to Select the Right Value of the Hyperparameter k

- **Cross-validation** is an experimental method for selecting the hyperparameter k
 - We will discuss it in greater detail in Part 6
- It consists of subdividing the initial data set D into three subsets:
 - training set, validation set, and test set
- We use the training set as the set containing the “prototypes”
- **Simple method:** we evaluate error E using different values of the k parameter with the validation set (more on this later)



The k-Nearest Neighbor (kNN) Method

- In the next slides, we see the theoretical concepts behind the k-NN method and how one can arrive at the approximation below:

$$\hat{P}(\omega_i | \mathbf{x}) = \frac{k_i}{k}$$

Density Estimation for the kNN Method

- The basic idea underlying many non-parametric methods can be illustrated as follows
- The probability P that a vector \mathbf{x} will fall in a region R of the feature space is:

$$P = \int_{\hat{A}} p(\mathbf{x}) d\mathbf{x}$$

- if R is a small region, P can be regarded as a smoothed or averaged version of the density function $p(\mathbf{x})$
- We can thus estimate this smoothed value of $p(\mathbf{x})$ with P

Density Estimation

- Suppose that n samples $\mathbf{x}_1, \dots, \mathbf{x}_n$ are drawn independently and identically distributed (i.i.d.) according to the probability law $p(\mathbf{x})$.
- If we know that \mathbf{k} samples out of these \mathbf{n} fall in R , then \mathbf{P} can be estimated as $\mathbf{P} = k/n$

- **Proof.** The probability that \mathbf{k} out of \mathbf{n} samples fall in R is given by the binomial law:

$$P_k = \binom{n}{k} P^k (1 - P)^{n-k}$$

- We know that the expected value of the binomial distribution for k is : $e(k) = nP$

Density Estimation

- If we consider the ratio k/n as argument of the binomial distribution, we can rewrite the expected value for k/n as:

$$e(k/n) = P$$


- We know that the variance of the binomial law is $\text{var}(k/n) = P(1-P)/n$
- If the number of samples n increases, the limit when it goes to infinity is:

$$n \rightarrow +\infty \rightarrow e(k/n) = P \text{ and } \text{var}(k/n) = 0$$

- Therefore, we can say that k/n is an asymptotically unbiased estimator of P
- We expect that the ratio k/n will be a very good estimate for the probability P , and hence for the smoothed density function if n is very large
 - Indeed, this estimate is especially accurate when n is very large (see next Figure)

Density Estimation: Binomial Distribution

- The plot shows the probability P_k of finding k samples in a region where the averaged probability is $P=0.7$, as a function of k/n
- Each curve is labeled by the total number of patterns n
- **For large n , such binomial distributions peak strongly at $k/n = P$ (here chosen to be 0.7)**



Density Estimation

- If we now assume that $p(\mathbf{x})$ is continuous and that **the region R is so small** that $p(\mathbf{x})$ does not vary appreciably within it, we can write:

$$\int_{\hat{A}} p(\mathbf{x}) d\mathbf{x} \approx p(\mathbf{x})V = P$$

- where \mathbf{x} is a point within R and V is the volume enclosed by R .
- Combining these two estimates, we get the following expression for $p(\mathbf{x})$:

$$\int_{\mathcal{R}} p(\mathbf{x}') d\mathbf{x}' \approx p(\mathbf{x})V \approx k/n \rightarrow p(\mathbf{x}) \approx \frac{k/n}{V}$$

Density Estimation

C. Bishop, *Pattern Recognition and Machine Learning*, pp. 121-122, 2006

$$p(\mathbf{x}) = \frac{k}{nV}$$

- The validity of the above estimate depends on two **contradictory** assumptions:
 1. the region R has to be small enough that the density is approximately constant over the region
 2. and yet sufficiently large (*in relation to the value of that density*) that the number k of points falling inside the region is sufficient for the binomial distribution to be sharply peaked
- We can exploit this result in two different ways
 - Fix k and determine the value of V from the data, which gives rise to the kNN method
 - Fix V and determine k from data, giving rise to kernel density estimators (Parzen windows)
- It can be shown that both the kNN density estimator and the kernel density estimator converge to the true probability density in the limit of infinite samples, provided that V shrinks suitably with n and k grows with n (Duda and Hart, 1973).

Density Estimation: Contradictory Assumptions

$$\int_{\mathfrak{R}} p(\mathbf{x}') d\mathbf{x}' @ p(\mathbf{x})V @ k/n \rightarrow p(\mathbf{x}) @ \frac{k/n}{V}$$

- **Bias:** If we fix the volume V and take more and more training samples, the ratio k/n will converge (in probability) as desired, but we have only obtained an estimate of the space-averaged value of $p(\mathbf{x})$:

$$\frac{P}{V} = \frac{\int_{\hat{A}} p(\mathbf{x}) d\mathbf{x}}{\hat{A}}$$

- **Variance:** if we want to obtain $p(\mathbf{x})$ rather than just a smoothed version of it, V has to tend to zero

Density Estimation

- From a practical standpoint, we note that the number of samples is always limited. Thus, the volume V can not be allowed to become arbitrarily small
 - One will have to accept a certain amount of *variance* in the ratio k/n and a certain amount of *bias*, i.e., averaging of the density $p(x)$
- From a theoretical standpoint, it is interesting to consider what happens in the limit of infinite samples
 - In practice, you have to ensure that, while the region shrinks (to reduce bias), the number of samples k within that region tends to infinity but at a lower rate than n (to reduce variance)

Density Estimation: Convergence

- To estimate the density at \mathbf{x} , we form a sequence of regions R_1, R_2, \dots, R_n containing \mathbf{x} while varying the total number of samples n
- Let V_n be the volume of R_n , k_n be the number of samples falling in R_n , and $p_n(\mathbf{x})$ be the n^{th} estimate for $p(\mathbf{x})$:

$$p_n(\mathbf{x}) = \frac{k_n/n}{V_n}$$

- If $p_n(\mathbf{x})$ is to converge to $p(\mathbf{x})$, three conditions appear to be required (as number of samples goes to infinity):

$$\lim_{n \rightarrow \infty} V_n = 0$$

$$\lim_{n \rightarrow \infty} k_n = \infty$$

$$\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$$

Density Estimation: Convergence

- The first condition assures us that the space averaged P/V will converge to $p(\mathbf{x})$, provided that the regions shrink uniformly and that $p(\cdot)$ is continuous at \mathbf{x} .
- The second condition, which only makes sense if $p(\mathbf{x}) \neq 0$, assures us that the frequency ratio will converge (in probability) to the probability P
- The third condition is necessary if $p_n(\mathbf{x})$ is to converge at all. It also says that although many samples will eventually fall within the small region R_n , they will form a negligibly small fraction of the total number of samples.

$$\lim_{n \rightarrow \infty} V_n = 0$$


$$\lim_{n \rightarrow \infty} k_n = \infty$$

$$\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$$

Density Estimation: Convergence

- There are two common ways of obtaining sequences of regions that satisfy these conditions
- One is to shrink an initial region by specifying the volume V_n as some function of n , such as $V_n = 1/\sqrt[n]{n}$. It then must be shown that the random variables k_n and k_n/n behave properly, or more to the point, that $p_n(x)$ converges to $p(x)$
 - This is the **Parzen-window / kernel density estimation (KDE)** method
- The second method is to specify k_n as some function of n , such as $k_n = \sqrt[n]{n}$. Here, the volume V_n is grown until it encloses k_n neighbors of x
 - This is the **kNN** method
- Both these methods converge in the asymptotic regime of infinite samples, but it is hard to give guarantees in the finite-sample regime


Density Estimation: KDE and kNN



k-Nearest Neighbor Method: Value of k


➤ kNN non-parametric method

- This method specifies k_n as some function of n , such as $k_n = \sqrt{n}$. Here, the volume V_n grows until it encloses k_n neighbors of x .
- **Key concept:** the region (volume) is specified considering the number k of the samples that fall into it.



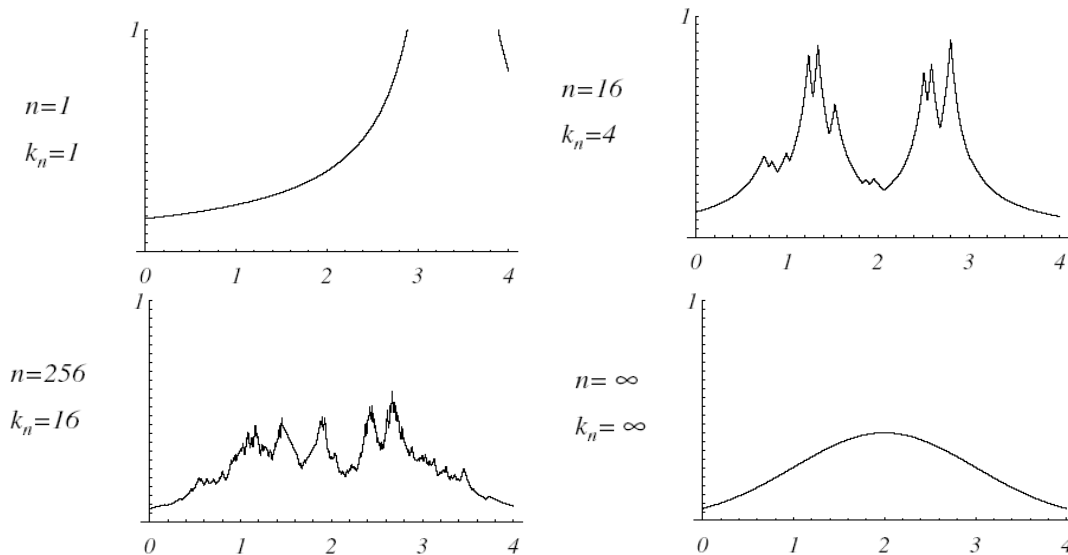
An Example of Density Estimation with kNN

- Eight points in one dimension and the kNN density estimates for $k=3$ and $k=5$.
- The discontinuities in the slopes in the estimates generally occur away from the positions of the points themselves



An Example of Density Estimation with kNN

- Several kNN estimates of a Gaussian distribution
- Note how the finite n estimates can be quite “spiky”



Kernel Density Estimation

- Recall that:

$$p(\mathbf{x}) = \frac{k}{nV}$$

- The validity of the above estimate depends on two **contradictory** assumptions:
 - the region R has to be small enough that the density is approximately constant over the region
 - and yet sufficiently large (*in relation to the value of that density*) that the number k of points falling inside the region is sufficient for the binomial distribution to be sharply peaked
- We can exploit this result in two different ways
 - Fix k and determine the value of V from the data, which gives rise to the kNN method
 - Fix V and determine k from data, giving rise to kernel density estimators (Parzen windows)**

Kernel Density Estimation

C. Bishop, *PRML*, pp. 123-124, 2006

- Let's consider R to be a small hypercube centered on the origin
 - This is referred to as a *kernel* function or Parzen window

$$k(\mathbf{u}) = \begin{cases} 1, & |u_i| \leq 1/2, \\ 0, & \text{otherwise} \end{cases} \quad i = 1, \dots, D,$$

- The quantity $k\left(\frac{\mathbf{x} - \mathbf{x}_n}{h}\right)$ will be one if the data point \mathbf{x}_n lies inside a cube of side h centered on \mathbf{x} , and zero otherwise
- The total number of data points lying inside this cube will therefore be

$$K = \sum_{n=1}^N k\left(\frac{\mathbf{x} - \mathbf{x}_n}{h}\right)$$

Kernel Density Estimation



- We can thus estimate $p(\mathbf{x})$ as

$$p(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N \frac{1}{h^D} k\left(\frac{\mathbf{x} - \mathbf{x}_n}{h}\right)$$

- If we consider a Gaussian kernel instead of a hypercube:

$$p(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N \frac{1}{(2\pi h^2)^{1/2}} \exp\left\{-\frac{\|\mathbf{x} - \mathbf{x}_n\|^2}{2h^2}\right\}$$

Kernel Density Estimation: 1D Examples



**We have discussed density estimation so far.
How about classification?**

From Density Estimation to Classification

- We can use the kNN or KDE estimators to get an estimate of the likelihood $p(\mathbf{x} | y)$
- Then estimate priors to compute the posterior probabilities, using the Bayes' theorem
- ... and decide for the class exhibiting the maximum support (MAP criterion)

The k-Nearest Neighbor (kNN) Method

- In the previous slides, we have shown that:

$$p_n(\mathbf{x}) = \frac{k_n/n}{V_n}$$

- It is easy to show that the kNN method estimates the posterior probabilities as

$$P(\omega_i|\mathbf{x}) = \frac{k_i}{k}$$

- where k_i is the number of nearest neighbors belonging to the class ω_i within the region R

The k-Nearest Neighbor (kNN) Method

- To arrive at this classification method, we fix k_n and n and allow for a variable V_n . Assuming Euclidean distance, let R be the region containing exactly k_n of the elements of the training set D . We know that the unconditional p.d.f. can be approximated as

$$p_n(\mathbf{x}) = \frac{k_n/n}{V_n}$$

- Denoting by k_i the number of elements in R from class ω_i , the class-conditional pdf for ω_i , $i=1, \dots, c$, can be approximated in R , as

$$p(\mathbf{x} / \omega_i) = \frac{k_i/n_i}{V_n}$$

The k-Nearest Neighbor (kNN) Method

- The posterior probabilities are obtained as

$$P_n(W_i | \mathbf{x}) = \frac{p_n(\mathbf{x} / W_i)P(W_i)}{p_n(x)} = \frac{\frac{k_i/n_i}{V} \frac{n_i}{n}}{\frac{k/n}{V}} = \frac{k_i}{k}$$


- The minimum error (Bayes) classifier using the approximations above will assign \mathbf{x} to the class with the highest posterior probability, i.e., the class most represented among the k nearest neighbors of \mathbf{x}
- The region R and the volume V are specific for each \mathbf{x}
 - The kNN classification rule, however, assigns the class label using only the numbers k_i , so the winning label does not depend on V

The k-Nearest Neighbor (kNN) Method

- kNN is Bayes-optimal when:

$$\lim_{n \rightarrow \infty} k_n = \infty$$

$$\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$$



- The error rate for the kNN rule for a two-category problem. Each curve is labeled by k; when $k=\infty$, the estimated probabilities match the true probabilities, and thus, the error rate is equal to the Bayes rate, i.e., $P = P_{\text{Bayes}}$

The Nearest-Neighbor Rule (1NN)


- Let be $D^n = \{x_1, \dots, x_n\}$ a training set of samples belonging to the “c” classes $\omega_1, \dots, \omega_c$ and be $\mathbf{x}' \in D^n$ the nearest sample to the unknown sample \mathbf{x}
- Decision rule: we assign \mathbf{x} to the class of \mathbf{x}'
- The nearest-neighbor rule is a sub-optimal procedure
 - that usually leads to an error rate greater than the minimum possible, the Bayes rate.
 - However, with infinite prototypes, the error rate is never worse than twice the Bayes rate.

The Nearest-Neighbor Rule (1NN)

- Let be ω' the true class of \mathbf{x}' . If “n” is very large, we can assume that \mathbf{x}' is very close to \mathbf{x} , so that $P(\omega' | \mathbf{x}') = P(\omega_i | \mathbf{x})$
 - Here, the nearest-neighbor decision rule is in good agreement with the MAP rule.
- In general, if we specify $\omega_m(\mathbf{x})$ by $P(W_m | \mathbf{x}) = \max_i P(W_i | \mathbf{x})$
- The MAP rule assigns the pattern to the class ω_m
 - If \mathbf{x}' has been assigned to the class j , then we should assume that $\omega_m(\mathbf{x}') = \omega_j$
- This rule partitions the feature space into cells that contain the closest points to a given training point \mathbf{x}'
 - All points in such a cell are thus labeled as the same class of the training point — a so-called **Voronoi tessellation** of the space (see next slide)

1NN and Voronoi Tessellation: A 2D Example

- In two dimensions, the nearest-neighbor algorithm **with $k=1$ (1-NN)** partitions the input space into Voronoi cells, each labeled by the category of the training point it contains




The Nearest-Neighbor Rule: Final Remarks

- If $P(\omega_m | \mathbf{x}) \cong 1$, the nearest-neighbor rule is close to the optimal rule, as it is very unlikely that the posterior probability changes abruptly.
 - When the minimum probability of error is small, the nearest-neighbor rule's probability of error is also small.
- If $P(\omega_m | \mathbf{x}) \cong 1/c$, the classes have the same probabilities, and the nearest-neighbor rule is likely suboptimal
 - In this case, the error probability is about $1-1/c$ for both methods.

Exercise (Laboratory)

- Consider the exercise solved with the Gaussian classifier in Part 2
- Solve it implementing a KDE classifier that estimates the likelihood of each class $p(x | y)$ with KDE
- Is it any different? Which one is more accurate?
- Test the two models on the moon dataset. *Which one is more accurate? Why?*



References

- Sections 4.1, 4.2, 4.3, 4.4, 4.5, Pattern Classification, R. O. Duda, P. E. Hart, D. G. Stork, John Wiley & Sons, 2000
- Chris Bishop, Machine Learning and Pattern Recognition, 2006.
- L. Kuncheva, Combining pattern classifiers, Wiley, 2004.